

# Python Pandas - Sorting

There are two kinds of sorting available in Pandas. They are –

- By label
- By Actual Value

Let us consider an example with an output.

```
import pandas as pd
import numpy as np

unsorted_df=pd.DataFrame(np.random.randn(10,2),index=[1,4,6,2,3,5,9,8,0,7],columns=['col2','col1'])
print unsorted_df
```

Its output is as follows –

	col2	col1
1	-2.063177	0.537527
4	0.142932	-0.684884
6	0.012667	-0.389340
2	-0.548797	1.848743
3	-1.044160	0.837381
5	0.385605	1.300185
9	1.031425	-1.002967
8	-0.407374	-0.435142
0	2.237453	-1.067139
7	-1.445831	-1.701035

In **unsorted\_df**, the **labels** and the **values** are unsorted. Let us see how these can be sorted.

## By Label

Using the **sort\_index()** method, by passing the axis arguments and the order of sorting, DataFrame can be sorted. By default, sorting is done on row labels in ascending order.

```
import pandas as pd
import numpy as np

unsorted_df =
pd.DataFrame(np.random.randn(10,2),index=[1,4,6,2,3,5,9,8,0,7],columns=['col2','col1'])

sorted_df=unsorted_df.sort_index()
print sorted_df
```

Its output is as follows –

	col2	col1
0	0.208464	0.627037
1	0.641004	0.331352
2	-0.038067	-0.464730
3	-0.638456	-0.021466
4	0.014646	-0.737438
5	-0.290761	-1.669827
6	-0.797303	-0.018737
7	0.525753	1.628921
8	-0.567031	0.775951
9	0.060724	-0.322425

## Order of Sorting

By passing the Boolean value to ascending parameter, the order of the sorting can be controlled. Let us consider the following example to understand the same.

```
import pandas as pd
import numpy as np

unsorted_df =
pd.DataFrame(np.random.randn(10,2), index=[1,4,6,2,3,5,9,8,0,7], co
lu
    mns = ['col2','col1'])

sorted_df = unsorted_df.sort_index(ascending=False)
print sorted_df
```

Its output is as follows –

	col2	col1
9	0.825697	0.374463
8	-1.699509	0.510373
7	-0.581378	0.622958
6	-0.202951	0.954300
5	-1.289321	-1.551250
4	1.302561	0.851385
3	-0.157915	-0.388659
2	-1.222295	0.166609
1	0.584890	-0.291048
0	0.668444	-0.061294

## Sort the Columns

By passing the axis argument with a value 0 or 1, the sorting can be done on the column labels. By default, axis=0, sort by row. Let us consider the following example to understand the same.

```
import pandas as pd
import numpy as np

unsorted_df =
pd.DataFrame(np.random.randn(10,2), index=[1,4,6,2,3,5,9,8,0,7], co
lu
```

```

mns = ['col2','col1'])

sorted_df=unsorted_df.sort_index(axis=1)

print sorted_df

```

Its output is as follows –

	col1	col2
1	-0.291048	0.584890
4	0.851385	1.302561
6	0.954300	-0.202951
2	0.166609	-1.222295
3	-0.388659	-0.157915
5	-1.551250	-1.289321
9	0.374463	0.825697
8	0.510373	-1.699509
0	-0.061294	0.668444
7	0.622958	-0.581378

## By Value

Like index sorting, **sort\_values()** is the method for sorting by values. It accepts a 'by' argument which will use the column name of the DataFrame with which the values are to be sorted.

```

import pandas as pd
import numpy as np

unsorted_df = pd.DataFrame({'col1':[2,1,1,1], 'col2':[1,3,2,4]})
sorted_df = unsorted_df.sort_values(by='col1')

print sorted_df

```

Its output is as follows –

	col1	col2
1	1	3
2	1	2
3	1	4
0	2	1

Observe, col1 values are sorted and the respective col2 value and row index will alter along with col1. Thus, they look unsorted.

'by' argument takes a list of column values.

```

import pandas as pd
import numpy as np

unsorted_df = pd.DataFrame({'col1':[2,1,1,1], 'col2':[1,3,2,4]})
sorted_df = unsorted_df.sort_values(by=['col1','col2'])

print sorted_df

```

Its output is as follows –

```
   col1 col2
2     1    2
1     1    3
3     1    4
0     2    1
```

## Sorting Algorithm

**sort\_values()** provides a provision to choose the algorithm from mergesort, heapsort and quicksort. Mergesort is the only stable algorithm.

```
import pandas as pd
import numpy as np

unsorted_df = pd.DataFrame({'col1':[2,1,1,1], 'col2':[1,3,2,4]})
sorted_df = unsorted_df.sort_values(by='col1' ,kind='mergesort')

print sorted_df
```

Its output is as follows –

```
   col1 col2
1     1    3
2     1    2
3     1    4
0     2    1
```

## Python Pandas - Working with Text Data

we will discuss the string operations with our basic Series/Index. In the subsequent chapters, we will learn how to apply these string functions on the DataFrame.

Pandas provides a set of string functions which make it easy to operate on string data. Most importantly, these functions ignore (or exclude) missing/NaN values.

Almost, all of these methods work with Python string functions (refer: <https://docs.python.org/3/library/stdtypes.html#string-methods>). So, convert the Series Object to String Object and then perform the operation.

Let us now see how each operation performs.

Sr.No	Function & Description
1	<b>lower()</b> Converts strings in the Series/Index to lower case.
2	<b>upper()</b> Converts strings in the Series/Index to upper case.

3	<b>len()</b> Computes String length().
4	<b>strip()</b> Helps strip whitespace(including newline) from each string in the Series/index from both the sides.
5	<b>split(' ')</b> Splits each string with the given pattern.
6	<b>cat(sep=' ')</b> Concatenates the series/index elements with given separator.
7	<b>get_dummies()</b> Returns the DataFrame with One-Hot Encoded values.
8	<b>contains(pattern)</b> Returns a Boolean value True for each element if the substring contains in the element, else False.
9	<b>replace(a,b)</b> Replaces the value <b>a</b> with the value <b>b</b> .
10	<b>repeat(value)</b> Repeats each element with specified number of times.
11	<b>count(pattern)</b> Returns count of appearance of pattern in each element.
12	<b>startswith(pattern)</b> Returns true if the element in the Series/Index starts with the pattern.
13	<b>endswith(pattern)</b> Returns true if the element in the Series/Index ends with the pattern.
14	<b>find(pattern)</b> Returns the first position of the first occurrence of the pattern.
15	<b>findall(pattern)</b>

	Returns a list of all occurrence of the pattern.
16	<b>swapcase</b> Swaps the case lower/upper.
17	<b>islower()</b> Checks whether all characters in each string in the Series/Index in lower case or not. Returns Boolean
18	<b>isupper()</b> Checks whether all characters in each string in the Series/Index in upper case or not. Returns Boolean.
19	<b>isnumeric()</b> Checks whether all characters in each string in the Series/Index are numeric. Returns Boolean.

Let us now create a Series and see how all the above functions work.

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t', np.nan,
               '1234', 'SteveSmith'])

print s
```

Its output is as follows –

```
0      Tom
1  William Rick
2      John
3   Alber@t
4        NaN
5      1234
6  Steve Smith
dtype: object
```

### lower()

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t', np.nan,
               '1234', 'SteveSmith'])

print s.str.lower()
```

Its output is as follows –

```
0      tom
1  william rick
2      john
```

```
3      alber@t
4      NaN
5      1234
6      steve smith
dtype: object
```

## upper()

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t', np.nan,
               '1234', 'SteveSmith'])

print s.str.upper()
```

Its output is as follows –

```
0      TOM
1  WILLIAM RICK
2      JOHN
3    ALBER@T
4      NaN
5    1234
6  STEVE SMITH
dtype: object
```

## len()

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t', np.nan,
               '1234', 'SteveSmith'])
print s.str.len()
```

Its output is as follows –

```
0      3.0
1     12.0
2      4.0
3      7.0
4      NaN
5      4.0
6     10.0
dtype: float64
```

## strip()

```
import pandas as pd
import numpy as np
s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])
print s
print ("After Stripping:")
print s.str.strip()
```

Its output is as follows –

```
0      Tom
```

```
1    William Rick
2           John
3       Alber@t
dtype: object
```

After Stripping:

```
0           Tom
1    William Rick
2           John
3       Alber@t
dtype: object
```

## split(pattern)

```
import pandas as pd
import numpy as np
s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])
print s
print ("Split Pattern:")
print s.str.split(' ')
```

Its output is as follows –

```
0           Tom
1    William Rick
2           John
3       Alber@t
dtype: object

Split Pattern:
0    [Tom, , , , , , , , , ]
1    [, , , , , William, Rick]
2    [John]
3    [Alber@t]
dtype: object
```

## cat(sep=pattern)

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print s.str.cat(sep='_')
```

Its output is as follows –

```
Tom _ William Rick_John_Alber@t
```

## get\_dummies()

```
import pandas as pd
import numpy as np

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print s.str.get_dummies()
```



Its output is as follows –

	William Rick	Alber@t	John	Tom
0	0	0	0	1
1	1	0	0	0
2	0	0	1	0
3	0	1	0	0

## contains ()

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print s.str.contains(' ')
```

Its output is as follows –

```
0    True
1    True
2   False
3   False
dtype: bool
```

## replace(a,b)

```
import pandas as pd
s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])
print s
print ("After replacing @ with $:")
print s.str.replace('@','$')
```

Its output is as follows –

```
0    Tom
1  William Rick
2    John
3  Alber@t
dtype: object

After replacing @ with $:
0    Tom
1  William Rick
2    John
3  Alber$t
dtype: object
```

## repeat(value)

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print s.str.repeat(2)
```

Its output is as follows –

0	Tom	Tom
1	William Rick	William Rick

```
2           JohnJohn
3           Alber@tAlber@t
dtype: object
```

## count(pattern)

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print ("The number of 'm's in each string:")
print s.str.count('m')
```

Its output is as follows –

```
The number of 'm's in each string:
0      1
1      1
2      0
3      0
```

## startswith(pattern)

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print ("Strings that start with 'T':")
print s.str.startswith('T')
```

Its output is as follows –

```
0    True
1    False
2    False
3    False
dtype: bool
```

## endswith(pattern)

```
import pandas as pd
s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])
print ("Strings that end with 't':")
print s.str.endswith('t')
```

Its output is as follows –

```
Strings that end with 't':
0    False
1    False
2    False
3     True
dtype: bool
```

## find(pattern)

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])
```

```
print s.str.find('e')
```

Its output is as follows –

```
0    -1
1    -1
2    -1
3     3
dtype: int64
```

"-1" indicates that there no such pattern available in the element.

## findall(pattern)

```
import pandas as pd

s = pd.Series(['Tom ', ' William Rick', 'John', 'Alber@t'])

print s.str.findall('e')
```

Its output is as follows –

```
0 []
1 []
2 []
3 [e]
dtype: object
```

Null list([ ]) indicates that there is no such pattern available in the element.

## swapcase()

```
import pandas as pd

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t'])

print s.str.swapcase()
```

Its output is as follows –

```
0    tOM
1  wILLIAM rICK
2   jOHN
3  aLBER@T
dtype: object
```

## islower()

```
import pandas as pd

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t'])

print s.str.islower()
```

Its output is as follows –

```
0    False
1    False
2    False
3    False
dtype: bool
```

## isupper()

```
import pandas as pd

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t'])

print s.str.isupper()
```

Its output is as follows –

```
0    False
1    False
2    False
3    False
dtype: bool
```

## isnumeric()

```
import pandas as pd

s = pd.Series(['Tom', 'William Rick', 'John', 'Alber@t'])

print s.str.isnumeric()
```

Its output is as follows –

```
0    False
1    False
2    False
3    False
dtype: bool
```

# Python Pandas - Options and Customization

Pandas provide API to customize some aspects of its behavior, display is being mostly used.

The API is composed of five relevant functions. They are –

- `get_option()`
- `set_option()`
- `reset_option()`
- `describe_option()`
- `option_context()`

Let us now understand how the functions operate.

## get\_option(param)

`get_option` takes a single parameter and returns the value as given in the output below –

## display.max\_rows

Displays the default number of value. Interpreter reads this value and displays the rows with this value as upper limit to display.

```
import pandas as pd
print pd.get_option("display.max_rows")
```

Its output is as follows –

```
60
```

## display.max\_columns

Displays the default number of value. Interpreter reads this value and displays the rows with this value as upper limit to display.

```
import pandas as pd
print pd.get_option("display.max_columns")
```

Its output is as follows –

```
20
```

Here, 60 and 20 are the default configuration parameter values.

## set\_option(param,value)

set\_option takes two arguments and sets the value to the parameter as shown below –

## display.max\_rows

Using **set\_option()**, we can change the default number of rows to be displayed.

```
import pandas as pd

pd.set_option("display.max_rows",80)

print pd.get_option("display.max_rows")
```

Its output is as follows –

```
80
```

## display.max\_columns

Using **set\_option()**, we can change the default number of rows to be displayed.

```
import pandas as pd

pd.set_option("display.max_columns",30)

print pd.get_option("display.max_columns")
```

Its output is as follows –

```
30
```

## reset\_option(param)

**reset\_option** takes an argument and sets the value back to the default value.

## display.max\_rows

Using `reset_option()`, we can change the value back to the default number of rows to be displayed.

```
import pandas as pd

pd.reset_option("display.max_rows")
print pd.get_option("display.max_rows")
```

Its output is as follows –

```
60
```

## describe\_option(param)

**describe\_option** prints the description of the argument.

## display.max\_rows

Using `reset_option()`, we can change the value back to the default number of rows to be displayed.

```
import pandas as pd
pd.describe_option("display.max_rows")
```

Its output is as follows –

```
display.max_rows : int
    If max_rows is exceeded, switch to truncate view. Depending on
    'large_repr', objects are either centrally truncated or
printed as
    a summary view. 'None' value means unlimited.

    In case python/IPython is running in a terminal and
`large_repr`
    equals 'truncate' this can be set to 0 and pandas will auto-
detect
    the height of the terminal and print a truncated object which
fits
    the screen height. The IPython notebook, IPython qtconsole, or
    IDLE do not run in a terminal and hence it is not possible to
do
    correct auto-detection.
[default: 60] [currently: 60]
```

## option\_context()

`option_context` context manager is used to set the option in **with statement** temporarily. Option values are restored automatically when you exit the **with block** –

## display.max\_rows

Using `option_context()`, we can set the value temporarily.

```
import pandas as pd
with pd.option_context("display.max_rows",10):
    print(pd.get_option("display.max_rows"))
print(pd.get_option("display.max_rows"))
```

Its **output** is as follows –

```
10
10
```

See, the difference between the first and the second print statements. The first statement prints the value set by **option\_context()** which is temporary within the **with context** itself. After the **with context**, the second print statement prints the configured value.

## Frequently used Parameters

Sr.No	Parameter & Description
1	<b>display.max_rows</b> Displays maximum number of rows to display
2	<b>2 display.max_columns</b> Displays maximum number of columns to display
3	<b>display.expand_frame_repr</b> Displays DataFrames to Stretch Pages
4	<b>display.max_colwidth</b> Displays maximum column width
5	<b>display.precision</b> Displays precision for decimal numbers

## Python Pandas - Indexing and Selecting Data

we will discuss how to slice and dice the data and generally get the subset of pandas object.

The Python and NumPy indexing operators "[ ]" and attribute operator "." provide quick and easy access to Pandas data structures across a wide range of use cases. However, since the type of the data to be accessed isn't known in advance, directly

using standard operators has some optimization limits. For production code, we recommend that you take advantage of the optimized pandas data access methods explained in this chapter.

Pandas now supports three types of Multi-axes indexing; the three types are mentioned in the following table –

Sr.No	Indexing & Description
1	<b>.loc()</b> Label based
2	<b>.iloc()</b> Integer based
3	<b>.ix()</b> Both Label and Integer based

## .loc()

Pandas provide various methods to have purely **label based indexing**. When slicing, the start bound is also included. Integers are valid labels, but they refer to the label and not the position.

**.loc()** has multiple access methods like –

- A single scalar label
- A list of labels
- A slice object
- A Boolean array

**loc** takes two single/list/range operator separated by ','. The first one indicates the row and the second one indicates columns.

## Example 1

```
#import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4),
index = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'], columns = ['A', 'B',
'C', 'D'])

#select all rows for a specific column
print df.loc[:, 'A']
```

Its **output** is as follows –



```
a    0.391548
b   -0.070649
c   -0.317212
d   -2.162406
e    2.202797
f    0.613709
g    1.050559
h    1.122680
Name: A, dtype: float64
```

## Example 2

```
# import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4),
index = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'], columns = ['A', 'B',
'C', 'D'])

# Select all rows for multiple columns, say list[]
print df.loc[:, ['A', 'C']]
```

Its output is as follows –

	A	C
a	0.391548	0.745623
b	-0.070649	1.620406
c	-0.317212	1.448365
d	-2.162406	-0.873557
e	2.202797	0.528067
f	0.613709	0.286414
g	1.050559	0.216526
h	1.122680	-1.621420

## Example 3

```
# import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4),
index = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'], columns = ['A', 'B',
'C', 'D'])

# Select few rows for multiple columns, say list[]
print df.loc[['a', 'b', 'f', 'h'], ['A', 'C']]
```

Its output is as follows –

	A	C
a	0.391548	0.745623
b	-0.070649	1.620406
f	0.613709	0.286414
h	1.122680	-1.621420

## Example 4

```
# import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4),
index = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'], columns = ['A', 'B',
'C', 'D'])

# Select range of rows for all columns
print df.loc['a':'h']
```

Its output is as follows –

	A	B	C	D
a	0.391548	-0.224297	0.745623	0.054301
b	-0.070649	-0.880130	1.620406	1.419743
c	-0.317212	-1.929698	1.448365	0.616899
d	-2.162406	0.614256	-0.873557	1.093958
e	2.202797	-2.315915	0.528067	0.612482
f	0.613709	-0.157674	0.286414	-0.500517
g	1.050559	-2.272099	0.216526	0.928449
h	1.122680	0.324368	-1.621420	-0.741470

## Example 5

```
# import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4),
index = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'], columns = ['A', 'B',
'C', 'D'])

# for getting values with a boolean array
print df.loc['a']>0
```

Its output is as follows –

```
A    False
B     True
C    False
D    False
Name: a, dtype: bool
```

## .iloc()

Pandas provide various methods in order to get purely integer based indexing. Like python and numpy, these are **0-based** indexing.

The various access methods are as follows –

- An Integer
- A list of integers
- A range of values

## Example 1

```
# import the pandas library and aliasing as pd
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

# select all rows for a specific column
print df.iloc[:4]
```

Its output is as follows –

	A	B	C	D
0	0.699435	0.256239	-1.270702	-0.645195
1	-0.685354	0.890791	-0.813012	0.631615
2	-0.783192	-0.531378	0.025070	0.230806
3	0.539042	-1.284314	0.826977	-0.026251

## Example 2

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

# Integer slicing
print df.iloc[:4]
print df.iloc[1:5, 2:4]
```

Its output is as follows –

	A	B	C	D
0	0.699435	0.256239	-1.270702	-0.645195
1	-0.685354	0.890791	-0.813012	0.631615
2	-0.783192	-0.531378	0.025070	0.230806
3	0.539042	-1.284314	0.826977	-0.026251

  

	C	D
1	-0.813012	0.631615
2	0.025070	0.230806
3	0.826977	-0.026251
4	1.423332	1.130568

## Example 3

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

# Slicing through list of values
print df.iloc[[1, 3, 5], [1, 3]]
print df.iloc[1:3, :]
print df.iloc[:, 1:3]
```

Its output is as follows –

	B	D
1	0.890791	0.631615
3	-1.284314	-0.026251
5	-0.512888	-0.518930

  

	A	B	C	D
1	-0.685354	0.890791	-0.813012	0.631615
2	-0.783192	-0.531378	0.025070	0.230806

  

	B	C
0	0.256239	-1.270702
1	0.890791	-0.813012
2	-0.531378	0.025070
3	-1.284314	0.826977
4	-0.460729	1.423332
5	-0.512888	0.581409
6	-1.204853	0.098060
7	-0.947857	0.641358

## .ix()

Besides pure label based and integer based, Pandas provides a hybrid method for selections and subsetting the object using the .ix() operator.

### Example 1

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

# Integer slicing
print df.ix[:4]
```

Its output is as follows –

	A	B	C	D
0	0.699435	0.256239	-1.270702	-0.645195
1	-0.685354	0.890791	-0.813012	0.631615
2	-0.783192	-0.531378	0.025070	0.230806
3	0.539042	-1.284314	0.826977	-0.026251

### Example 2

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

# Index slicing
print df.ix[:, 'A']
```

Its output is as follows –

0	0.699435
---	----------

```
1 -0.685354
2 -0.783192
3  0.539042
4 -1.044209
5 -1.415411
6  1.062095
7  0.994204
Name: A, dtype: float64
```

## Use of Notations

Getting values from the Pandas object with Multi-axes indexing uses the following notation –

Object	Indexers	Return Type
Series	s.loc[indexer]	Scalar value
DataFrame	df.loc[row_index,col_index]	Series object
Panel	p.loc[item_index,major_index, minor_index]	p.loc[item_index,major_index, minor_index]

**Note – .iloc() & .ix()** applies the same indexing options and Return value.

Let us now see how each operation can be performed on the DataFrame object. We will use the basic indexing operator '[' –

### Example 1

```
import pandas as pd
import numpy as np
df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])
print df['A']
```

Its output is as follows –

```
0 -0.478893
1  0.391931
2  0.336825
3 -1.055102
4 -0.165218
5 -0.328641
6  0.567721
7 -0.759399
Name: A, dtype: float64
```

**Note –** We can pass a list of values to [ ] to select those columns.

### Example 2

```
import pandas as pd
```

```
import numpy as np
df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])

print df[['A', 'B']]
```

Its output is as follows –

	A	B
0	-0.478893	-0.606311
1	0.391931	-0.949025
2	0.336825	0.093717
3	-1.055102	-0.012944
4	-0.165218	1.550310
5	-0.328641	-0.226363
6	0.567721	-0.312585
7	-0.759399	-0.372696

### Example 3

```
import pandas as pd
import numpy as np
df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])
print df[2:2]
```

Its output is as follows –

```
Columns: [A, B, C, D]
Index: []
```

## Attribute Access

Columns can be selected using the attribute operator '.'.

### Example

```
import pandas as pd
import numpy as np
df = pd.DataFrame(np.random.randn(8, 4), columns = ['A', 'B', 'C', 'D'])
```

`print df.A` Statistical methods help in the understanding and analyzing the behavior of data. We will now learn a few statistical functions, which we can apply on Pandas objects.

## Percent\_change

Series, DataFrames and Panel, all have the function **pct\_change()**. This function compares every element with its prior element and computes the change percentage.

```
import pandas as pd
import numpy as np
s = pd.Series([1,2,3,4,5,4])
print s.pct_change()
```

```
df = pd.DataFrame(np.random.randn(5, 2))
print df.pct_change()
```

Its output is as follows –

```
0      NaN
1    1.000000
2    0.500000
3    0.333333
4    0.250000
5   -0.200000
dtype: float64

      0      1
0      NaN      NaN
1  -15.151902   0.174730
2   -0.746374  -1.449088
3   -3.582229  -3.165836
4    15.601150  -1.860434
```

By default, the **pct\_change()** operates on columns; if you want to apply the same row wise, then use **axis=1()** argument.

## Covariance

Covariance is applied on series data. The Series object has a method **cov** to compute covariance between series objects. NA will be excluded automatically.

### Cov Series

```
import pandas as pd
import numpy as np
s1 = pd.Series(np.random.randn(10))
s2 = pd.Series(np.random.randn(10))
print s1.cov(s2)
```

Its output is as follows –

```
-0.12978405324
```

Covariance method when applied on a DataFrame, computes **cov** between all the columns.

```
import pandas as pd
import numpy as np
frame = pd.DataFrame(np.random.randn(10, 5), columns=['a', 'b', 'c', 'd', 'e'])
print frame['a'].cov(frame['b'])
print frame.cov()
```

Its output is as follows –

```
-0.58312921152741437

      a      b      c      d      e
a  1.780628 -0.583129 -0.185575  0.003679 -0.136558
b -0.583129  1.297011  0.136530 -0.523719  0.251064
```

c	-0.185575	0.136530	0.915227	-0.053881	-0.058926
d	0.003679	-0.523719	-0.053881	1.521426	-0.487694
e	-0.136558	0.251064	-0.058926	-0.487694	0.960761

**Note** – Observe the **cov** between **a** and **b** column in the first statement and the same is the value returned by cov on DataFrame.

## Correlation

Correlation shows the linear relationship between any two array of values (series). There are multiple methods to compute the correlation like pearson(default), spearman and kendall.

```
import pandas as pd
import numpy as np
frame = pd.DataFrame(np.random.randn(10, 5), columns=['a', 'b', 'c', 'd', 'e'])

print frame['a'].corr(frame['b'])
print frame.corr()
```

Its output is as follows –

```
-0.383712785514
```

	a	b	c	d	e
a	1.000000	-0.383713	-0.145368	0.002235	-0.104405
b	-0.383713	1.000000	0.125311	-0.372821	0.224908
c	-0.145368	0.125311	1.000000	-0.045661	-0.062840
d	0.002235	-0.372821	-0.045661	1.000000	-0.403380
e	-0.104405	0.224908	-0.062840	-0.403380	1.000000

If any non-numeric column is present in the DataFrame, it is excluded automatically.

## Data Ranking

Data Ranking produces ranking for each element in the array of elements. In case of ties, assigns the mean rank.

```
import pandas as pd
import numpy as np

s = pd.Series(np.random.randn(5), index=list('abcde'))
s['d'] = s['b'] # so there's a tie
print s.rank()
```

Its output is as follows –

```
a  1.0
b  3.5
c  2.0
d  3.5
e  5.0
dtype: float64
```

Rank optionally takes a parameter ascending which by default is true; when false, data is reverse-ranked, with larger values assigned a smaller rank.



Rank supports different tie-breaking methods, specified with the method parameter –

- **average** – average rank of tied group
- **min** – lowest rank in the group
- **max** – highest rank in the group
- **first** – ranks assigned in the order they appear in the array

Its output is as follows –

```
0    -0.478893
1     0.391931
2     0.336825
3    -1.055102
4    -0.165218
5    -0.328641
6     0.567721
7    -0.759399
Name: A, dtype: float64
```

## Python Pandas - Statistical Functions

Statistical methods help in the understanding and analyzing the behavior of data. We will now learn a few statistical functions, which we can apply on Pandas objects.

### Percent\_change

Series, DataFrames and Panel, all have the function **pct\_change()**. This function compares every element with its prior element and computes the change percentage.

```
import pandas as pd
import numpy as np
s = pd.Series([1,2,3,4,5,4])
print s.pct_change()

df = pd.DataFrame(np.random.randn(5, 2))
print df.pct_change()
```

Its output is as follows –

```
0      NaN
1    1.000000
2    0.500000
3    0.333333
4    0.250000
5   -0.200000
dtype: float64
```

	0	1
0	NaN	NaN
1	-15.151902	0.174730
2	-0.746374	-1.449088
3	-3.582229	-3.165836
4	15.601150	-1.860434

By default, the **pct\_change()** operates on columns; if you want to apply the same row wise, then use **axis=1()** argument.

## Covariance

Covariance is applied on series data. The Series object has a method **cov** to compute covariance between series objects. NA will be excluded automatically.

### Cov Series

```
import pandas as pd
import numpy as np
s1 = pd.Series(np.random.randn(10))
s2 = pd.Series(np.random.randn(10))
print s1.cov(s2)
```

Its output is as follows –

```
-0.12978405324
```

Covariance method when applied on a DataFrame, computes **cov** between all the columns.

```
import pandas as pd
import numpy as np
frame = pd.DataFrame(np.random.randn(10, 5), columns=['a', 'b', 'c', 'd', 'e'])
print frame['a'].cov(frame['b'])
print frame.cov()
```

Its output is as follows –

```
-0.58312921152741437
```

	a	b	c	d	e
a	1.780628	-0.583129	-0.185575	0.003679	-0.136558
b	-0.583129	1.297011	0.136530	-0.523719	0.251064
c	-0.185575	0.136530	0.915227	-0.053881	-0.058926
d	0.003679	-0.523719	-0.053881	1.521426	-0.487694
e	-0.136558	0.251064	-0.058926	-0.487694	0.960761

**Note** – Observe the **cov** between **a** and **b** column in the first statement and the same is the value returned by **cov** on DataFrame.

## Correlation

Correlation shows the linear relationship between any two array of values (series). There are multiple methods to compute the correlation like **pearson**(default), **spearman** and **kendall**.

```
import pandas as pd
import numpy as np
frame = pd.DataFrame(np.random.randn(10, 5), columns=['a', 'b', 'c', 'd', 'e'])

print frame['a'].corr(frame['b'])
print frame.corr()
```

Its output is as follows –

```
-0.383712785514
```

	a	b	c	d	e
a	1.000000	-0.383713	-0.145368	0.002235	-0.104405
b	-0.383713	1.000000	0.125311	-0.372821	0.224908
c	-0.145368	0.125311	1.000000	-0.045661	-0.062840
d	0.002235	-0.372821	-0.045661	1.000000	-0.403380
e	-0.104405	0.224908	-0.062840	-0.403380	1.000000

If any non-numeric column is present in the DataFrame, it is excluded automatically.

## Data Ranking

Data Ranking produces ranking for each element in the array of elements. In case of ties, assigns the mean rank.

```
import pandas as pd
import numpy as np

s = pd.Series(np.random.randn(5), index=list('abcde'))
s['d'] = s['b'] # so there's a tie
print s.rank()
```

Its output is as follows –

```
a  1.0
b  3.5
c  2.0
d  3.5
e  5.0
dtype: float64
```

Rank optionally takes a parameter ascending which by default is true; when false, data is reverse-ranked, with larger values assigned a smaller rank.

Rank supports different tie-breaking methods, specified with the method parameter –

- **average** – average rank of tied group
- **min** – lowest rank in the group
- **max** – highest rank in the group
- **first** – ranks assigned in the order they appear in the array

# Python Pandas - Window Functions

For working on numerical data, Pandas provide few variants like rolling, expanding and exponentially moving weights for window statistics. Among these are **sum**, **mean**, **median**, **variance**, **covariance**, **correlation**, etc.

We will now learn how each of these can be applied on DataFrame objects.

## .rolling() Function

This function can be applied on a series of data. Specify the **window=n** argument and apply the appropriate statistical function on top of it.

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(10, 4),
                  index = pd.date_range('1/1/2000', periods=10),
                  columns = ['A', 'B', 'C', 'D'])
print df.rolling(window=3).mean()
```

Its output is as follows –

	A	B	C	D
2000-01-01	NaN	NaN	NaN	NaN
2000-01-02	NaN	NaN	NaN	NaN
2000-01-03	0.434553	-0.667940	-1.051718	-0.826452
2000-01-04	0.628267	-0.047040	-0.287467	-0.161110
2000-01-05	0.398233	0.003517	0.099126	-0.405565
2000-01-06	0.641798	0.656184	-0.322728	0.428015
2000-01-07	0.188403	0.010913	-0.708645	0.160932
2000-01-08	0.188043	-0.253039	-0.818125	-0.108485
2000-01-09	0.682819	-0.606846	-0.178411	-0.404127
2000-01-10	0.688583	0.127786	0.513832	-1.067156

**Note** – Since the window size is 3, for first two elements there are nulls and from third the value will be the average of the **n**, **n-1** and **n-2** elements. Thus we can also apply various functions as mentioned above.

## .expanding() Function

This function can be applied on a series of data. Specify the **min\_periods=n** argument and apply the appropriate statistical function on top of it.

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(10, 4),
                  index = pd.date_range('1/1/2000', periods=10),
                  columns = ['A', 'B', 'C', 'D'])
print df.expanding(min_periods=3).mean()
```

Its output is as follows –

	A	B	C	D
2000-01-01	NaN	NaN	NaN	NaN
2000-01-02	NaN	NaN	NaN	NaN
2000-01-03	0.434553	-0.667940	-1.051718	-0.826452
2000-01-04	0.743328	-0.198015	-0.852462	-0.262547
2000-01-05	0.614776	-0.205649	-0.583641	-0.303254
2000-01-06	0.538175	-0.005878	-0.687223	-0.199219
2000-01-07	0.505503	-0.108475	-0.790826	-0.081056
2000-01-08	0.454751	-0.223420	-0.671572	-0.230215
2000-01-09	0.586390	-0.206201	-0.517619	-0.267521
2000-01-10	0.560427	-0.037597	-0.399429	-0.376886

## .ewm() Function

**ewm** is applied on a series of data. Specify any of the **com**, **span**, **halflife** argument and apply the appropriate statistical function on top of it. It assigns the weights exponentially.

```
import pandas as pd
import numpy as np

df = pd.DataFrame(np.random.randn(10, 4),
                  index = pd.date_range('1/1/2000', periods=10),
                  columns = ['A', 'B', 'C', 'D'])
print df.ewm(com=0.5).mean()
```

Its output is as follows –

	A	B	C	D
2000-01-01	1.088512	-0.650942	-2.547450	-0.566858
2000-01-02	0.865131	-0.453626	-1.137961	0.058747
2000-01-03	-0.132245	-0.807671	-0.308308	-1.491002
2000-01-04	1.084036	0.555444	-0.272119	0.480111
2000-01-05	0.425682	0.025511	0.239162	-0.153290
2000-01-06	0.245094	0.671373	-0.725025	0.163310
2000-01-07	0.288030	-0.259337	-1.183515	0.473191
2000-01-08	0.162317	-0.771884	-0.285564	-0.692001
2000-01-09	1.147156	-0.302900	0.380851	-0.607976
2000-01-10	0.600216	0.885614	0.569808	-1.110113

Window functions are majorly used in finding the trends within the data graphically by smoothing the curve. If there is lot of variation in the everyday data and a lot of data points are available, then taking the samples and plotting is one method and applying the window computations and plotting the graph on the results is another method. By these methods, we can smooth the curve or the trend.